

Data Copilot.

Automating the first day of a new client engagement, without hiding what the AI did.

PRESENTED BY

Zahra Moghaddasi

SECTOR

Data & Analytics Consulting

COMPANY SIZE

Small

Every new client starts the same way.

Chleo runs a small data analytics consultancy. A raw export lands on a consultant's desk, and someone has to manually work through it before any billable analysis can start. That doesn't scale, and Chleo is scared that handing this to “the AI” means she can no longer explain to a client what it did.

- 01 Figure out what's actually in the export, and how clean it is
- 02 Work out how the tables relate to each other
- 03 Decide on a sensible reporting model
- 04 Pick the KPIs that matter to this specific client
- 05 Write the first-draft “here's what your data tells us” report

Three use cases, one pipeline.

01 / QUALITY

Data quality & schema profiling

Nulls, duplicates, outliers, and key relationships, confirmed by real value overlap, not just column names.

02 / MODEL + KPI

AI data model & KPI generator

Proposes a fact/dimension model and a KPI list with SQL, then actually runs it against the data.

03 / INSIGHT

Business insight engine

Turns the results into a plain-language client report, plus forward-looking data science ideas.

The client onboarding report.

IF THE DEMO DOESN'T COOPERATE, HERE ARE THE REAL NUMBERS, FROM A REAL RUN
FIGURES IN BRL, SEP 2016-OCT 2018 · ORDER VALUE EXCLUDES SHIPPING

TOTAL REVENUE

R\$ 13,591,643.70

AVG. ORDER VALUE

R\$ 137.75

AVG. REVIEW SCORE

4.09 / 5

TOP CATEGORY (BY REVENUE)

Health & Beauty

“Most orders are paid by credit card: 76,505 transactions.”

“261,831 duplicate rows in geolocation could affect location-based insights.”

The pipeline, and what LangSmith makes visible.

EVERY AI DECISION, LOGGED AND REVIEWABLE. NOT A BLACK BOX

☐	✓	Name	Input	Output	Error	Start Time	Latency	Dataset	Annotation ...	Tokens	Cost	First Token	Tags	Metadata	Feedback
☐	✓	ChatOpenAI	user: You are writing a sh...	chat.completion		8/31/2026, 12:35:25 AM	🕒 2.06s			1,601	\$0.0003297			LANGSMITH_TRACIN	
☐	✓	ChatOpenAI	user: You are a senior dat...	chat.completion		8/31/2026, 12:35:19 AM	🕒 5.47s			4,517	\$0.0007272			Is_temperature: 0.2	

Two traced calls per run: the model & KPI generation step (4,517 tokens, 5.47s), and the insight-writing step (1,601 tokens, 2.06s). Full prompt and response, both reviewable.

What this costs, and how fast it moves.

BUILD ITEM	HRS	COST
Quality & schema profiling	8	760 €
Model & KPI generation	10	950 €
Insight generation & dashboard	6	570 €
LangSmith & review workflow	3	285 €
Testing on client-like data	5	475 €
Total build	32	3,040 €

95 €/hr blended rate · Germany market · ~48 € per client onboarding
after build, vs. 380-760 € manual today

1 wk	Discovery scope with 1-2 pilot clients
2-3 wks	Pilot build harden against real client shapes
2-4 wks	Pilot run alongside the manual process
1 wk	Review decide go / no-go
~6-9 weeks to a go / no-go decision	

What I'm asking for.

THE ASK

Greenlight a short pilot: run this on 1-2 more real client-shaped datasets over the next few weeks, alongside the existing manual process, before committing further.