

Data Copilot, proven and compliant.

Three real client shapes, real scale, and the legal groundwork done before any real client data touches it.

PRESENTED BY

Zahra Moghaddasi

SECTOR

Data & Analytics Consulting

THIS ROUND

MVP + compliance + ROI

Every new client starts the same way.

Chleo runs a small data analytics consultancy. A raw export lands on a consultant's desk, and someone has to manually work through it before any billable analysis can start. That doesn't scale, and Chleo is scared that handing this to “the AI” means she can no longer explain to a client what it did.

- 01 Figure out what's actually in the export, and how clean it is
- 02 Work out how the tables relate to each other
- 03 Decide on a sensible reporting model
- 04 Pick the KPIs that matter to this specific client
- 05 Write the first-draft “here's what your data tells us” report

Three client shapes.

Round 1 proved the idea on one clean dataset.

01 / CLEAN

Olist

Brazilian E-Commerce Public Dataset by Olist

Relational e-commerce data. 1.5 million rows.

02 / MESSY, AT SCALE

CFPB

Consumer Financial Protection Bureau,
Complaint Database

One flat table, free-text heavy. Full 17.4 million row dataset, not a slice.

03 / SEMI-STRUCTURED

Open Food Facts

Open Food Facts (OFF), collaborative product database

Genuinely nested JSON. The AI proposes the relational tables itself.

New this round: a local PII and casing check runs before anything reaches OpenAI or a client, see next.

The onboarding report, live.

CFPB ROWS PROFILED

17,456,743

PROFILING TIME

~90 sec

KPIs COMPUTED

5 / 5

REAL COST, THIS RUN

\$0.0014

“The PII scan flagged personal names in free text before anything reached OpenAI.”

The pipeline, and what LangSmith makes visible.

EVERY AI DECISION, LOGGED AND REVIEWABLE. NOT A BLACK BOX

✓	Name	Input	Output	Error	Start Time	Latency	Dataset	Annotation ...	Tokens	Cost
✓	ChatOpenAI	user: You are writing a short business-...	chatcmpl-EKJ0DHzp6z7...		9/4/2026, 9:51:53 AM	🕒 2.19s			1.581K	\$0.0003276
✓	ChatOpenAI	user: You are a senior data consultant ...	chatcmpl-EKJ08BcmeTts...		9/4/2026, 9:51:47 AM	🕒 4.91s			4.856K	\$0.001020
✓	ChatOpenAI	user: Here is a field catalog for a client'...	chatcmpl-EK4xETIFWbz3...		9/3/2026, 6:51:51 PM	🕒 33.17s			7.921K	\$0.003535
✓	ChatOpenAI	user: Here is a field catalog for a client'...	chatcmpl-EK4wr3S6tn4...		9/3/2026, 6:51:28 PM	🕒 22.88s			7.145K	\$0.002604
✓	ChatOpenAI	user: Here is a field catalog for a client'...	chatcmpl-EK4wWTb8bL...		9/3/2026, 6:51:06 PM	🕒 21.95s			7.162K	\$0.002624

Four steps, real numbers from this round: profile, ask the AI for a model + KPIs, run the KPI SQL for real (self-correcting up to 5 times on failure), then write the plain-language insights. Full prompt and response for every AI call, and the real token cost measured straight from the API, not estimated.

Is this legal to run?

EU AI ACT

Minimal risk

- Not Annex III: no hiring, credit, benefits, or law-enforcement decisions
- B2B analytics on a client's own data, not a decision about a person
- The consultancy is a deployer, not a GPAI provider (that's OpenAI)
- Article 50(2) marking obligation doesn't apply: mandatory human review before publication is exactly its carve-out

GDPR

Client = controller

- Consultancy is the processor, acting on the client's instructions only
- Personal data redacted locally, never sent to OpenAI raw
- Tested for real: 56-63.5% of CFPB free-text narratives had a detected name, even after their own redaction
- Mandatory human review before any report reaches a client, the control that anchors both frameworks

What could still go wrong.

3 of 8 risks tracked in the full risk matrix, the ones that most need a human in the loop.

REGULATORY · L3 × I5

PII hidden in free text reaches OpenAI

Local scan and redaction before anything leaves the machine. Verified: over half of CFPB's narratives had a detectable name.

REGULATORY · L3 × I4

The PII check misses something

A first pass, not a guarantee, especially on non-English text. Mandatory human review before delivery stays the real control.

ETHICAL · L2 × I5

Over-reliance on AI output replaces the consultant's own judgment

Positioned as a first draft only, never sent straight to a client. Enforced as a mandatory review gate through the pilot phase.

Full matrix, all 8 risks (regulatory, technical, ethical, operational), likelihoods, impacts, and mitigations:
`roi_risk_assessment.md`

What this costs, and what it's worth.

Real LLM cost, measured per run	\$0.0014
Consultant review (30 min, mandatory)	48 €
Manual onboarding today (4-8 hrs)	380-760 €
Net saving per onboarding	~522 €

AI cost is negligible, measured directly from real API usage. The real cost is the 30-minute human review, which stays mandatory by design.

12 MONTHS

ROI113%

36 MONTHS

ROI370%

Recommended next steps.

DONE

POC

This capstone: MVP tested on 3 real datasets

NEXT

Pilot

2-3 real clients, mandatory human review gate stays permanent

LATER

Full deployment

Standard part of every new-client onboarding

THE ASK

Greenlight a pilot: 2-3 real clients, run alongside the existing manual process, compared directly on time and quality before we go further.

■ CLOSE

Thank you.

Questions?